

Name: Zhenjiang Tian

BlazerID: ztian

Course: CS 532

Assignment: HW4

## 1. Design and Implementation

### 1.1 Overall Architecture

The program uses one parent process and one child process created by fork().

A unidirectional pipe is used for communication: the parent writes integers to the pipe, while the child reads them.

- Parent process:

- Creates the pipe.
- Forks the child process.
- Closes the read-end of the pipe.
- Creates 10 producer threads.
- Waits for all producer threads to finish.
- Closes the write-end of the pipe.
- Sends SIGUSR1 to the child to indicate that all numbers have been generated.

- Child process:

- Installs a SIGUSR1 signal handler.
- Closes the write-end of the pipe.
- Waits using pause() until the SIGUSR1 signal is received.
- Creates 20 consumer threads to read from the pipe.
- Waits for all consumer threads to finish.
- Computes the average of the 20 partial sums and prints the result.

### 1.2 Producer Threads

The parent process creates NUM\_PRODUCERS = 10 threads. Each producer thread:

- Generates NUM\_PER\_PRODUCER = 500 unique random integers in the range [0, 1000].  
A local array (used[1001]) is maintained per thread to guarantee uniqueness within that thread.
- Uses rand\_r() with a per-thread seed to avoid data races in random generation.
- Before calling write() on the pipe, the thread acquires a mutex (pipe\_mutex).  
This ensures that at most one producer writes to the pipe at a time and prevents interleaving of partial writes.
- After write() completes, the thread releases the mutex.
- Periodically prints progress (every 50 numbers) and a final completion message.  
A separate mutex (print\_mutex) is used to keep console output readable.

This satisfies the requirements of using threads, generating random numbers, and applying synchronization to avoid race conditions and data corruption.

### 1.3 Consumer Threads

The child process creates `NUM_CONSUMERS = 20` threads after receiving `SIGUSR1`.

Each consumer thread:

- Reads `NUM_PER_CONSUMER = 250` integers from the pipe using `read()`.
- Accumulates these values into a local long long sum.
- Stores its result in a global array `consumer_sums[tid]`.
- Prints its own completion message and sum (also protected by `print_mutex`).

When all consumer threads have completed, the child process:

- Adds up all 20 entries of `consumer_sums`.
- Computes the average = `total / NUM_CONSUMERS`.
- Prints the final average to standard output, which is redirected to `result.txt`.

#### 1.4 Signal Mechanism (Graduate Requirement)

To meet the graduate requirement of signaling:

- The child process installs a signal handler for `SIGUSR1`:
    - The handler sets a global flag `start_reading = 1`.
  - Before creating any consumer threads, the child calls `pause()` in a loop:
    - while (`!start_reading`) `pause();`
  - The parent process sends `SIGUSR1` to the child using:
    - `kill(child_pid, SIGUSR1);`
- but only after:
- All producer threads have finished.
  - The write-end of the pipe has been closed.

This design guarantees that:

- The child does not start reading from the pipe until all numbers have been generated.
- The signal serves exactly as the notification mechanism required by the assignment.

#### 1.5 Synchronization and Data Integrity

- `pipe_mutex`: protects `write()` operations from concurrent producers.
- `print_mutex`: protects `printf()` output from multiple threads.
- The pipe itself guarantees ordered delivery of the written integers.
- On macOS, POSIX unnamed semaphores are not available, so mutexes are used instead to provide mutual exclusion around critical sections.

## 2. AI Tool Usage

### 2.1 Tool Used

- ChatGPT

### 2.2 How the AI Tool Helped

The AI tool was used for debugging.

### 2.3 Example Prompts Used

```
I wrote a C code, but the system prompts the following problem  
at runtime. I really don't know why. Can you tell me the reason  
and how I should modify it (base) ztian@Zhenjiangs-MacBook-  
Pro hw4 % gcc -Wall -Wextra -O2 -std=c11 hw4_ztian.c -lpthread  
-o hw4 hw4_ztian.c:134:13: warning: 'sem_init' is deprecated [-  
Wdeprecated-declarations] 134 | if (sem_init(&pipe_sem, 0, 1) !=  
0) { | ^  
/Library/Developer/CommandLineTools/SDKs/MacOSX.sdk/usr/in-  
clude/sys/semaphore.h:55:42: note: 'sem_init' has been explicitly  
marked deprecated here 55 | int sem_init(sem_t *, int, unsigned  
int) __deprecated; | ^  
/Library/Developer/CommandLineTools/SDKs/MacOSX.sdk/usr/in-  
clude/sys/cdefs.h:214:40: note: expanded from macro  
'__deprecated' 214 | #define __deprecated  
__attribute__((__deprecated__)) | ^ hw4_ztian.c:159:9: warning:  
'sem_destroy' is deprecated [-Wdeprecated-declarations] 159 |  
sem_destroy(&pipe_sem); | ^  
/Library/Developer/CommandLineTools/SDKs/MacOSX.sdk/usr/in-  
clude/sys/semaphore.h:53:26: note: 'sem_destroy' has been  
explicitly marked deprecated here 53 | int sem_destroy(sem_t *)  
__deprecated; | ^  
/Library/Developer/CommandLineTools/SDKs/MacOSX.sdk/usr/in-  
clude/sys/cdefs.h:214:40: note: expanded from macro  
'__deprecated' 214 | #define __deprecated  
__attribute__((__deprecated__)) | ^ 2 warnings generated. (base)  
ztian@Zhenjiangs-MacBook-Pro hw4 % ./hw4 > result.txt  
sem_init: Function not implemented
```

Thought for 24s >

This isn't your fault — it's macOS being weird. 😊

This isn't your fault — it's macOS being weird. 😊

Your code is basically fine for Linux, but **macOS does not support `sem_init` / `sem_destroy` for unnamed semaphores**, so at runtime it fails with:

text

Copy code

```
sem_init: Function not implemented
```

Let's break it down and then I'll show you how to fix it.

## 1. What the warning means

When you compile:

bash

Copy code

```
gcc -Wall -Wextra -O2 -std=c11 hw4_ztian.c -lpthread -o hw4
```

you see:

text

Copy code

```
warning: 'sem_init' is deprecated
warning: 'sem_destroy' is deprecated
```

On macOS, the header `/usr/include/sys/semaphore.h` marks `sem_init` and `sem_destroy` as `__deprecated`. That's the compiler telling you: "these functions are old and not really supported anymore".



Then at runtime, your program calls:

```
C

if (sem_init(&pipe_sem, 0, 1) != 0) {
    perror("sem_init");
    exit(1);
}
```

 Copy code

and `sem_init` returns `-1` with `errno = ENOSYS`, so `perror` prints:

text

 Copy code

```
sem_init: Function not implemented
```

So: **on macOS, unnamed POSIX semaphores (`sem_t`, `sem_init`, `sem_destroy`) simply don't work.**  
They're stubs.

On Linux they work fine, which is why your professor's code or slides probably use them.

---

## 2. What you can do instead

You have three realistic options:

### Option A: Run the same code on Linux (no code change)

If you want to keep using `sem_t` exactly as in your homework:

- Compile and run on:
  - Your course's Linux server, **or**
  - Ubuntu in a VM, **or**
  - WSL on Windows (if you ever use Windows),

On Linux, the same code:



```
sem_t pipe_sem;

if (sem_init(&pipe_sem, 0, 1) != 0) {
    perror("sem_init");
    exit(1);
}

...
sem_destroy(&pipe_sem);
```

 Copy code

will work and not print "Function not implemented".

If your professor will grade on a Linux machine anyway, this is the *safest* option for homework correctness.

---

### Option B: Replace the semaphore with a `pthread_mutex_t` (for macOS)

If you're just using the semaphore as a **lock** (to protect some shared data or a pipe), a mutex does the same job for threads.

1. At the top of your file, replace the semaphore with a mutex:

```
c

#include <pthread.h>
// #include <semaphore.h> // you can remove this if you no longer use sem_t

pthread_mutex_t pipe_lock = PTHREAD_MUTEX_INITIALIZER;
```

 Copy code

2. Remove your `sem_init` and `sem_destroy` calls (e.g. at lines 134 and 159 in your code).
3. Wherever you had:



**3.** Wherever you had:

```
C Copy code
```

```
sem_wait(&pipe_sem);
/* critical section */
sem_post(&pipe_sem);
```

change to:

```
C Copy code
```

```
pthread_mutex_lock(&pipe_lock);
/* critical section */
pthread_mutex_unlock(&pipe_lock);
```

**4.** If you want explicit init/destroy (optional, since we used `PTHREAD_MUTEX_INITIALIZER`):

```
C Copy code
```

```
// in main(), before creating threads:
if (pthread_mutex_init(&pipe_lock, NULL) != 0) {
    perror("pthread_mutex_init");
    exit(1);
}

// at the very end of main():
pthread_mutex_destroy(&pipe_lock);
```



Now there will be **no** `sem_init` **call**, so no “Function not implemented”.

This change is usually enough for your homework if the assignment only needs mutual exclusion and not counting semantics between processes.

So I just followed the steps of option B and modified my code, and then it worked well.